

# Sparse & irregular processing on GPUs using functional languages

**Adam Harries**

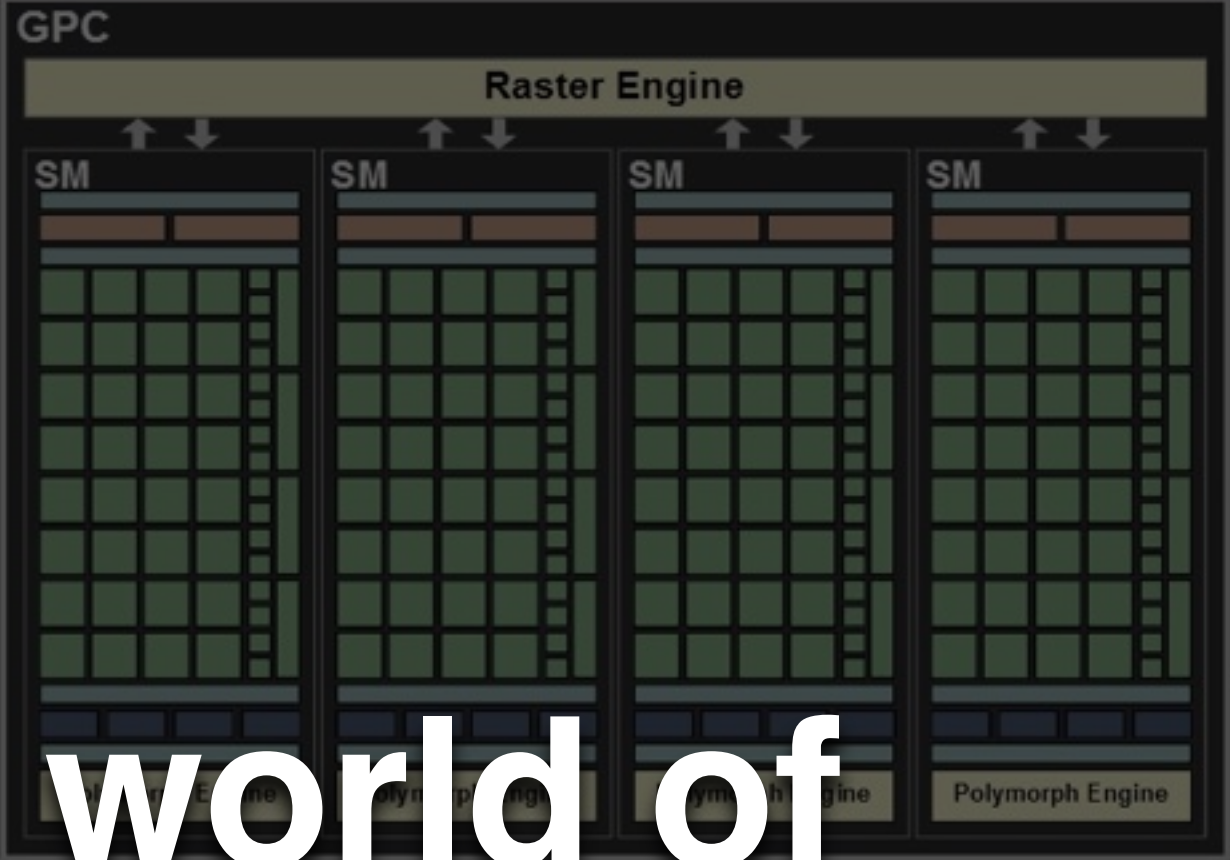
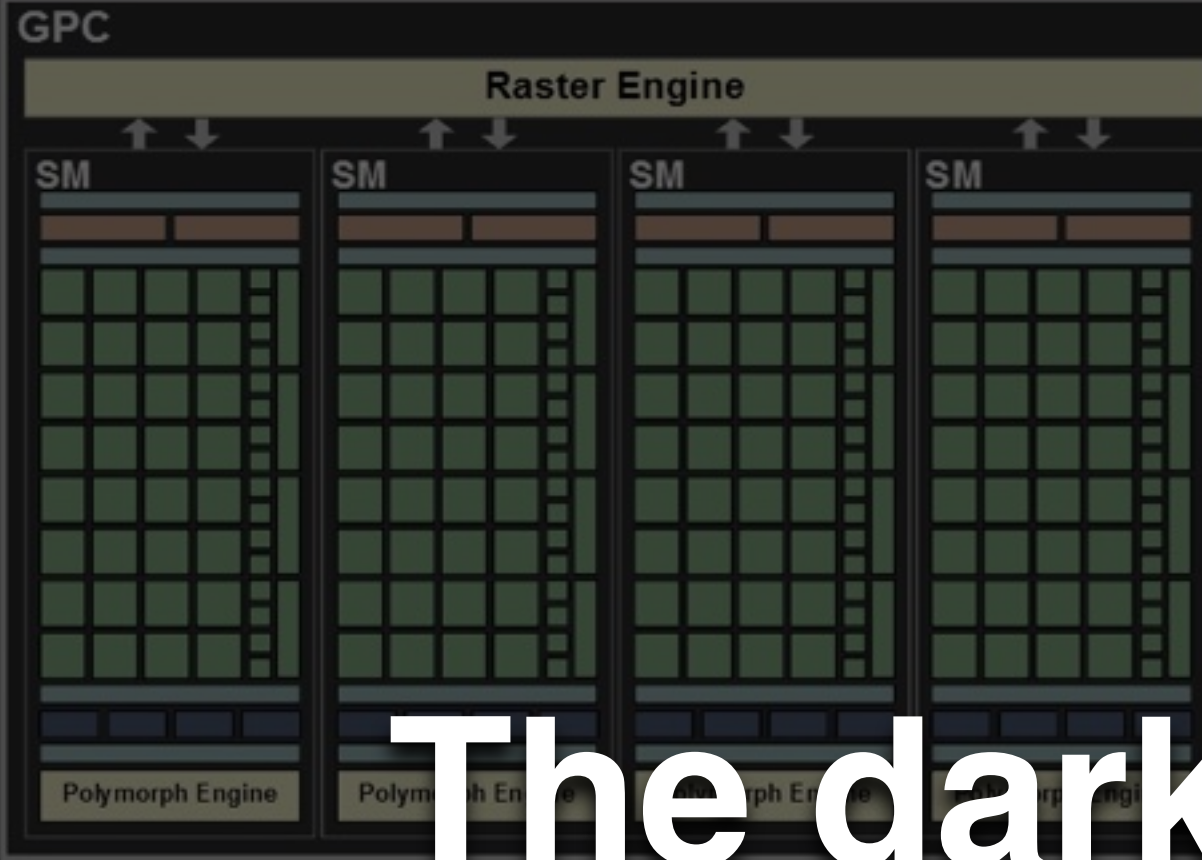
School of Informatics, University of Edinburgh



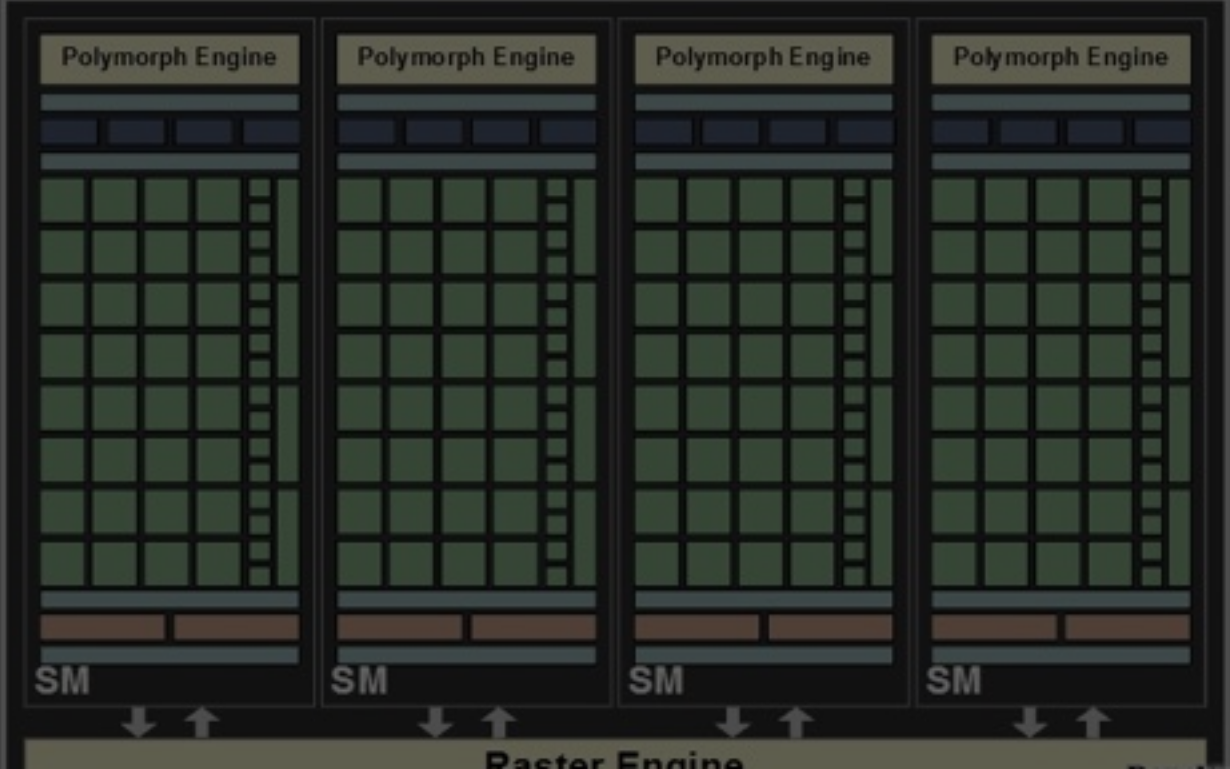
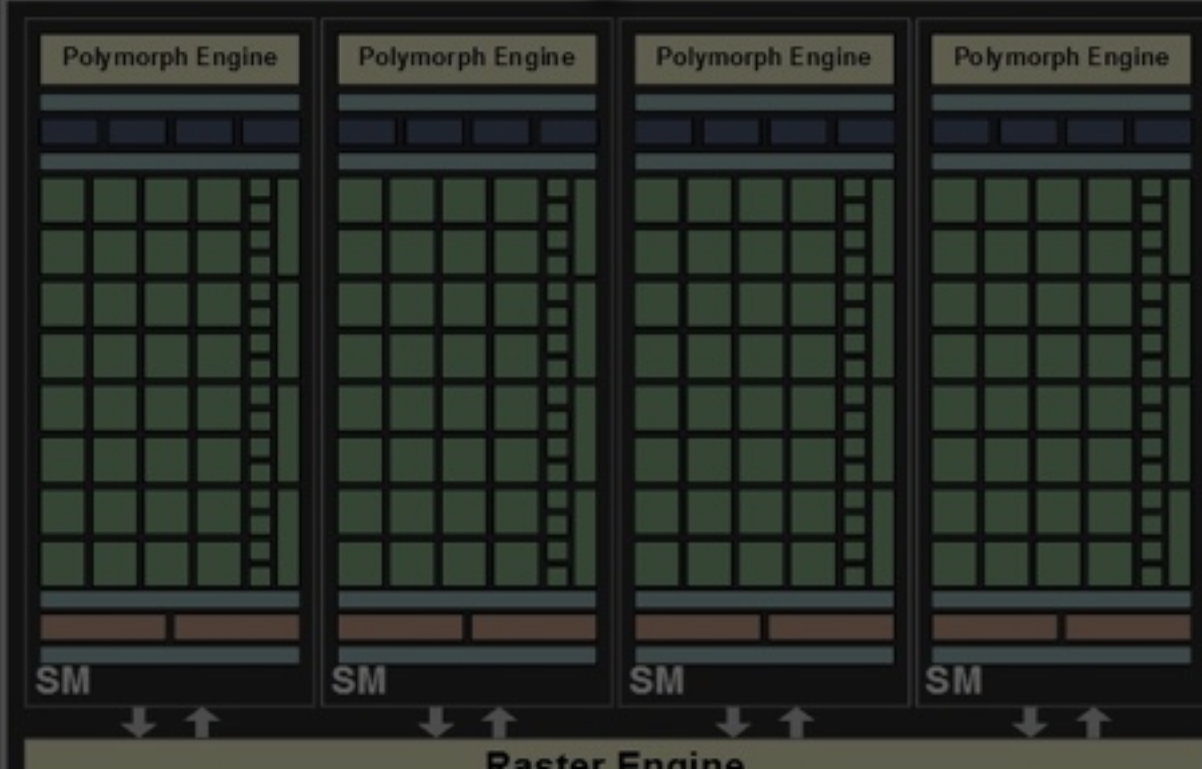
THE UNIVERSITY of EDINBURGH  
**informatics**

**EPSRC** Centre for Doctoral Training in  
**Pervasive Parallelism**

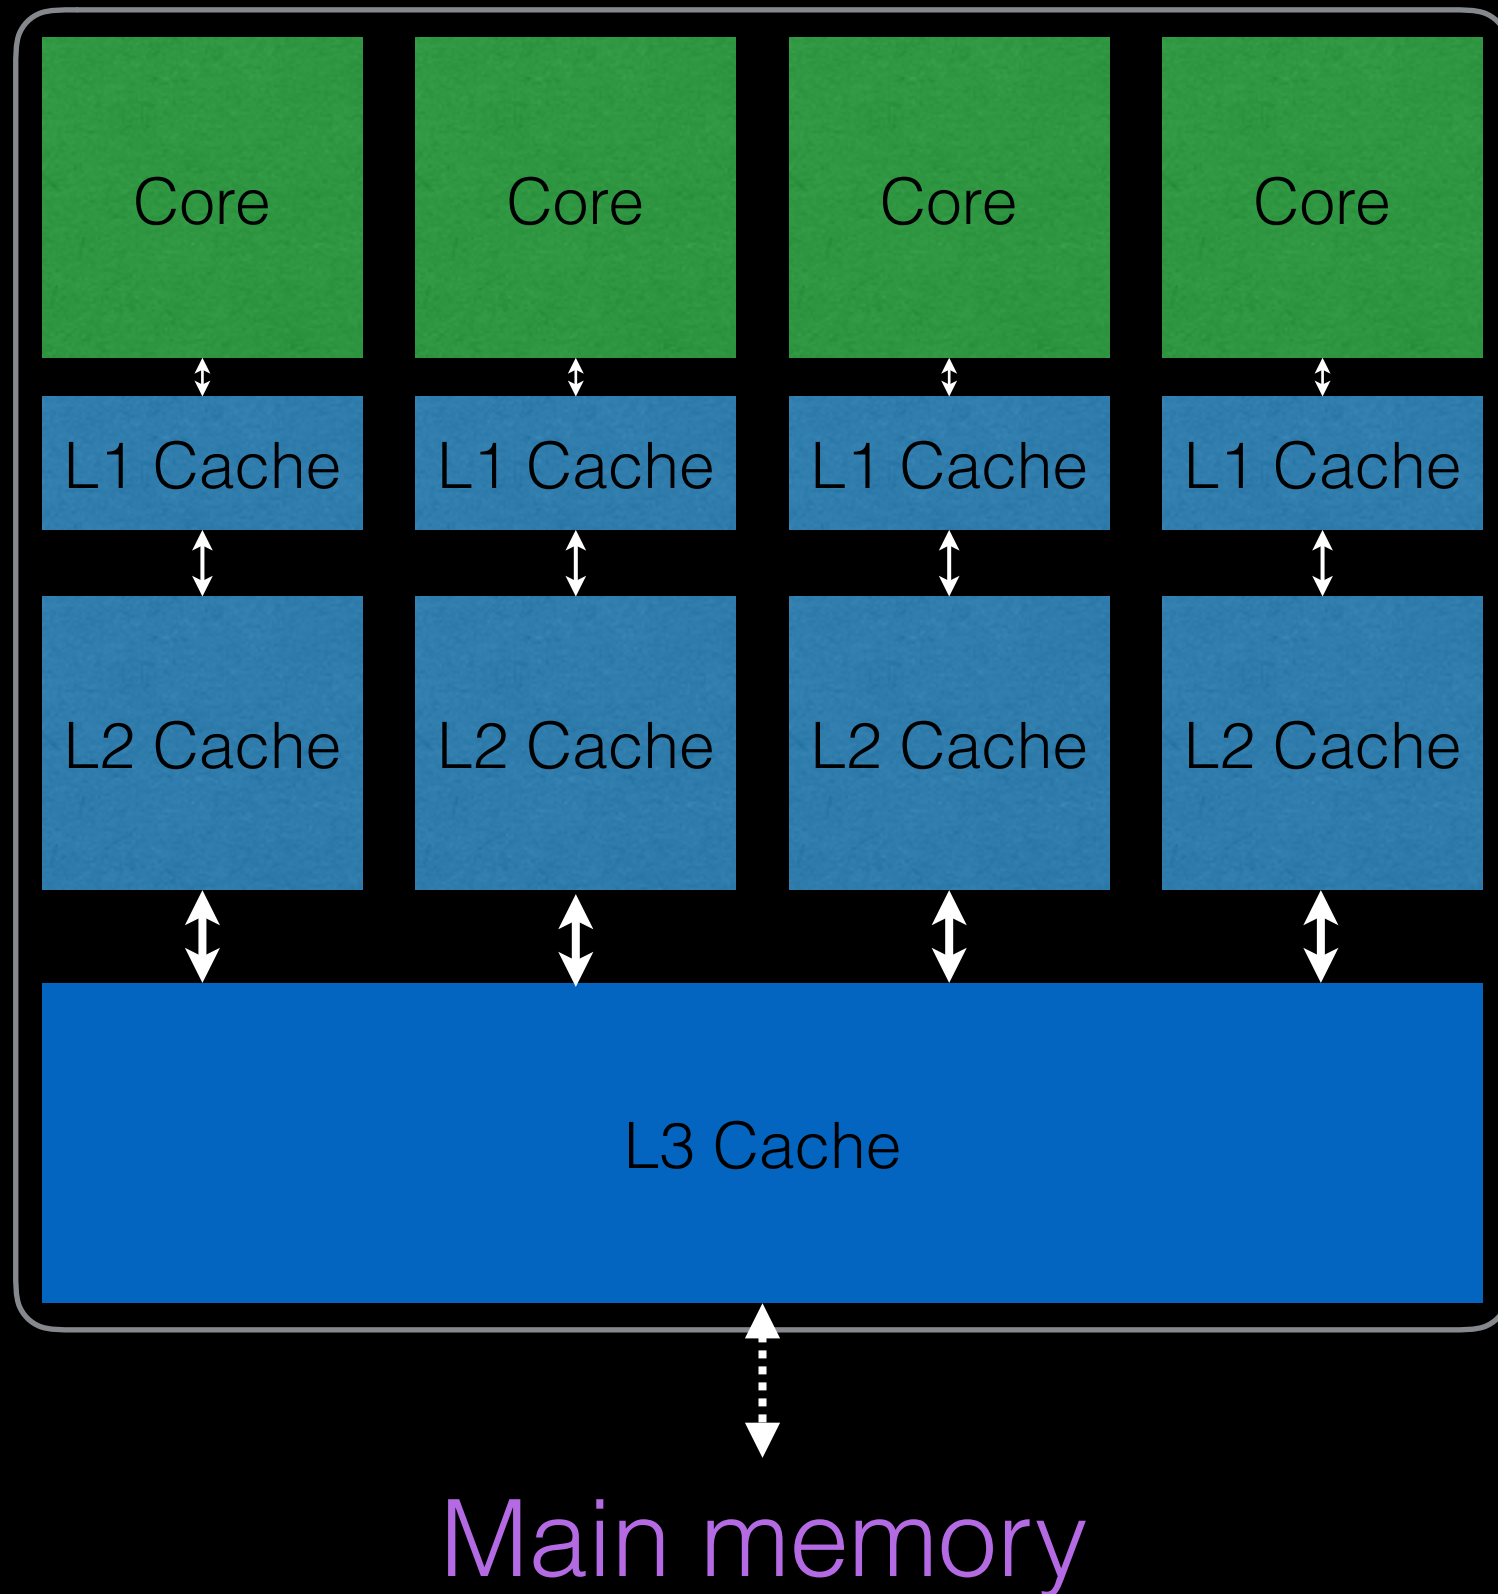
**EPSRC**  
Engineering and Physical Sciences  
Research Council



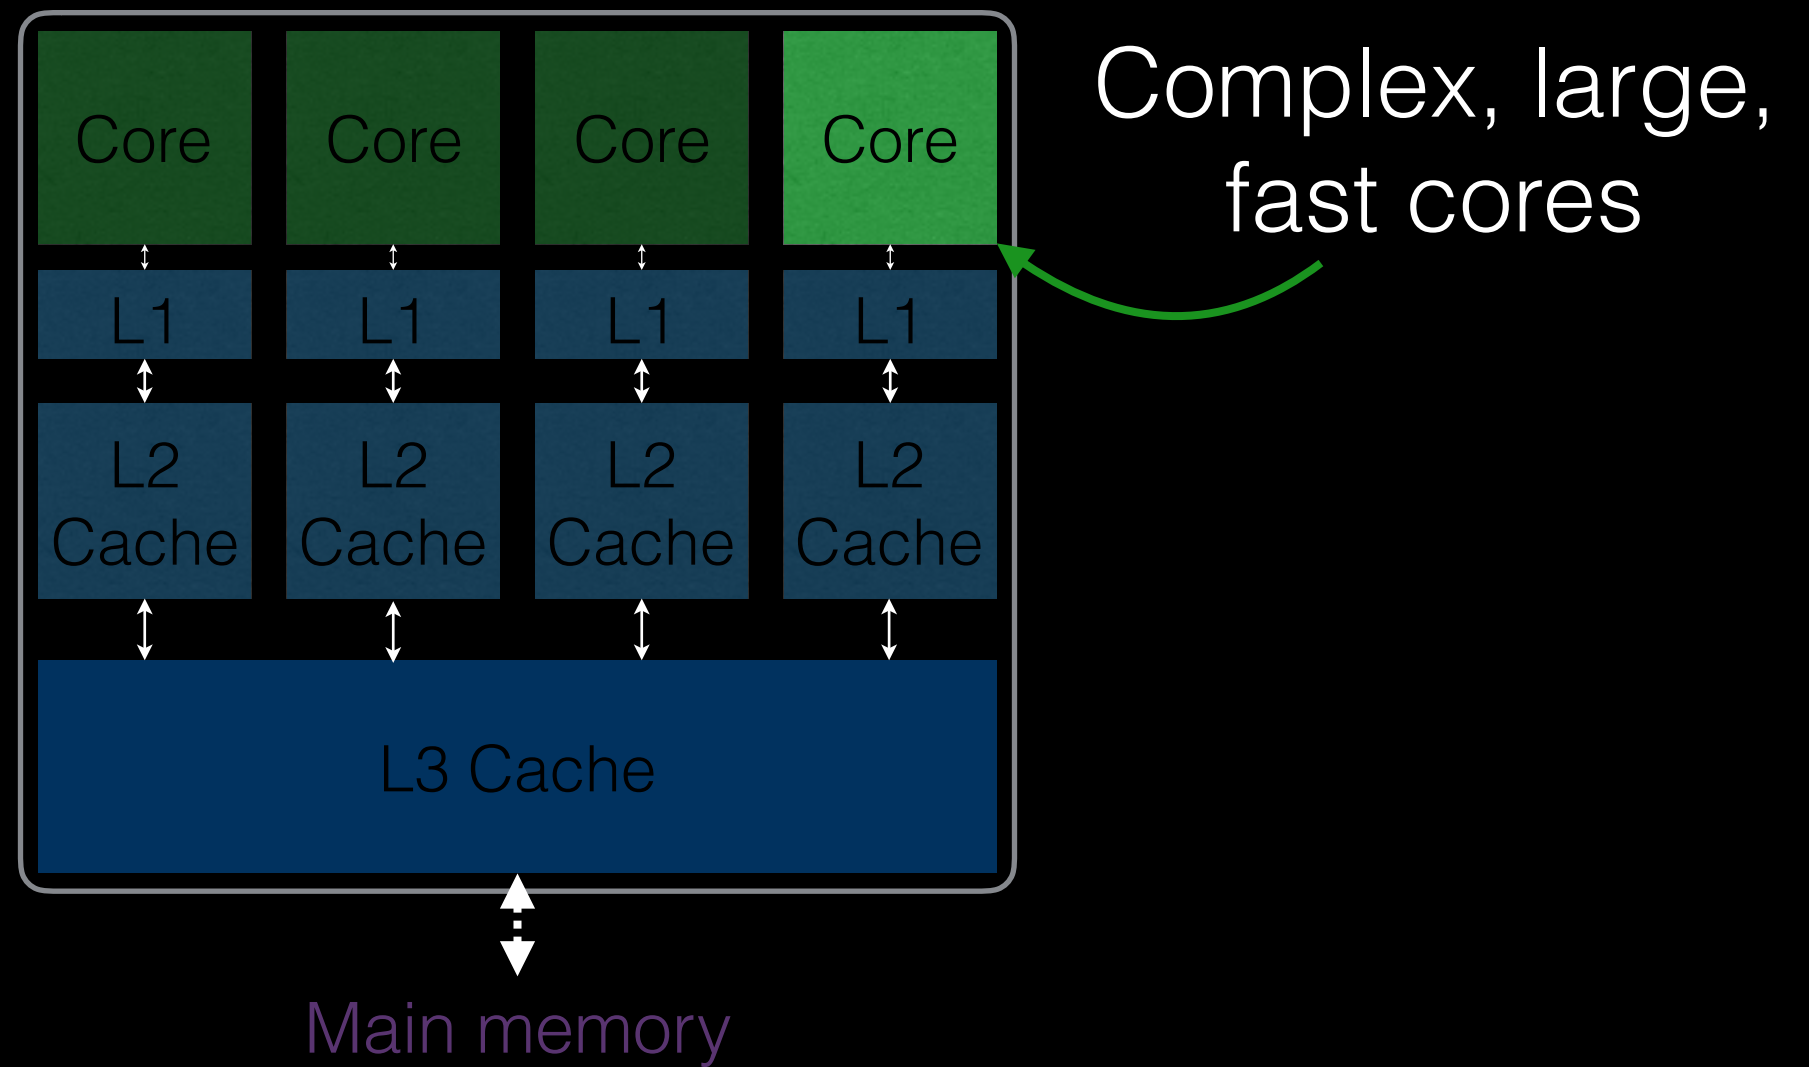
# The dark world of computer architecture



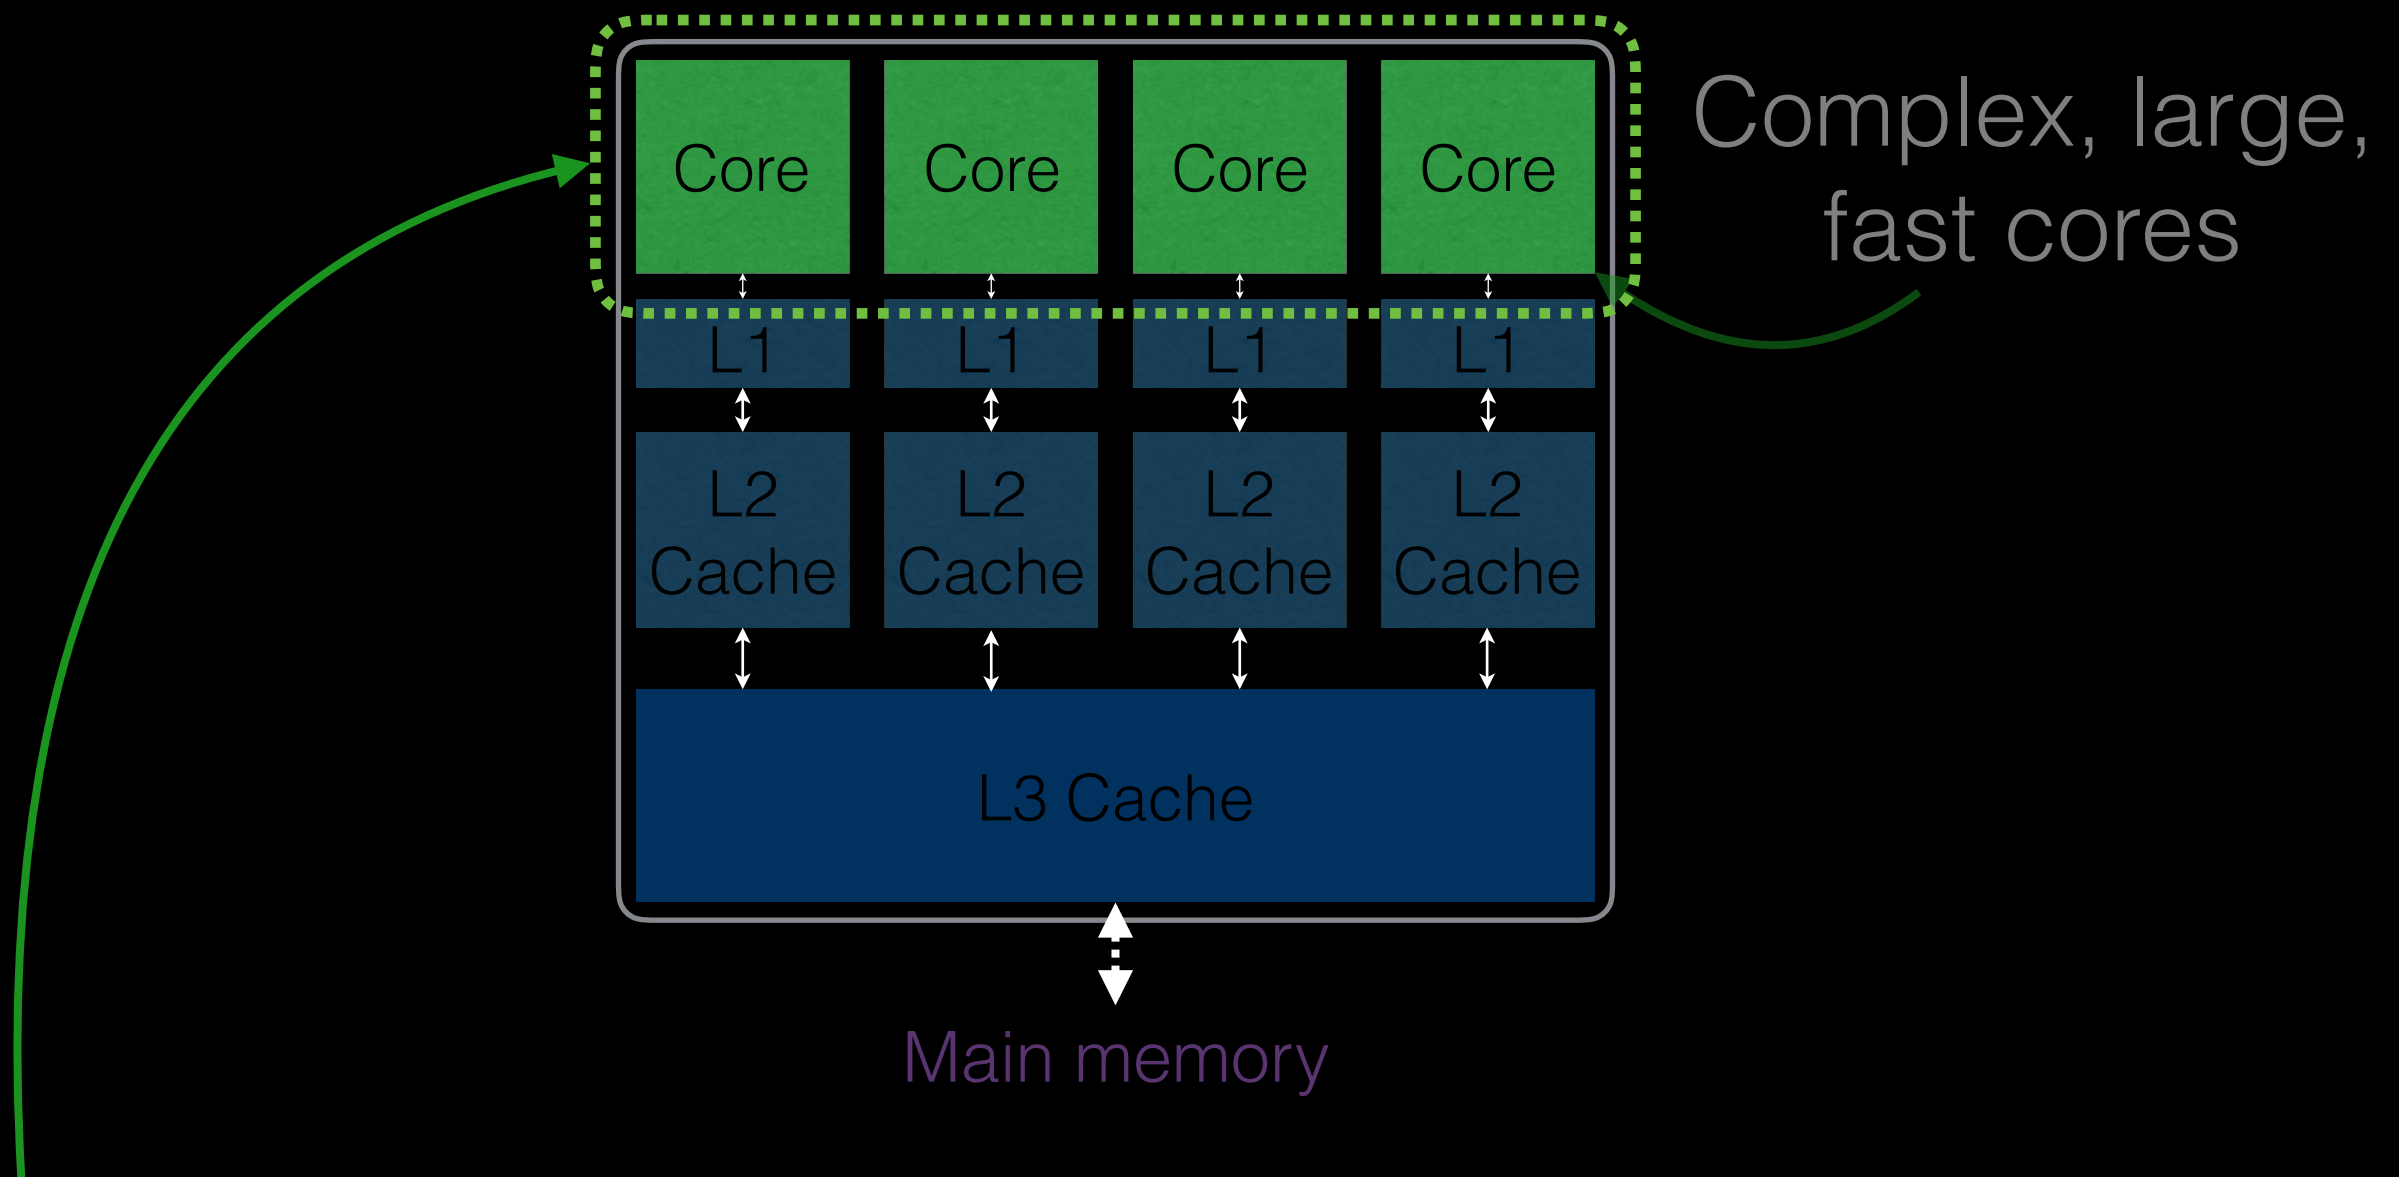
# Current CPUs



# Current CPUs




# Current CPUs



Coarse parallelism: **low thread count** and **heavyweight threads**

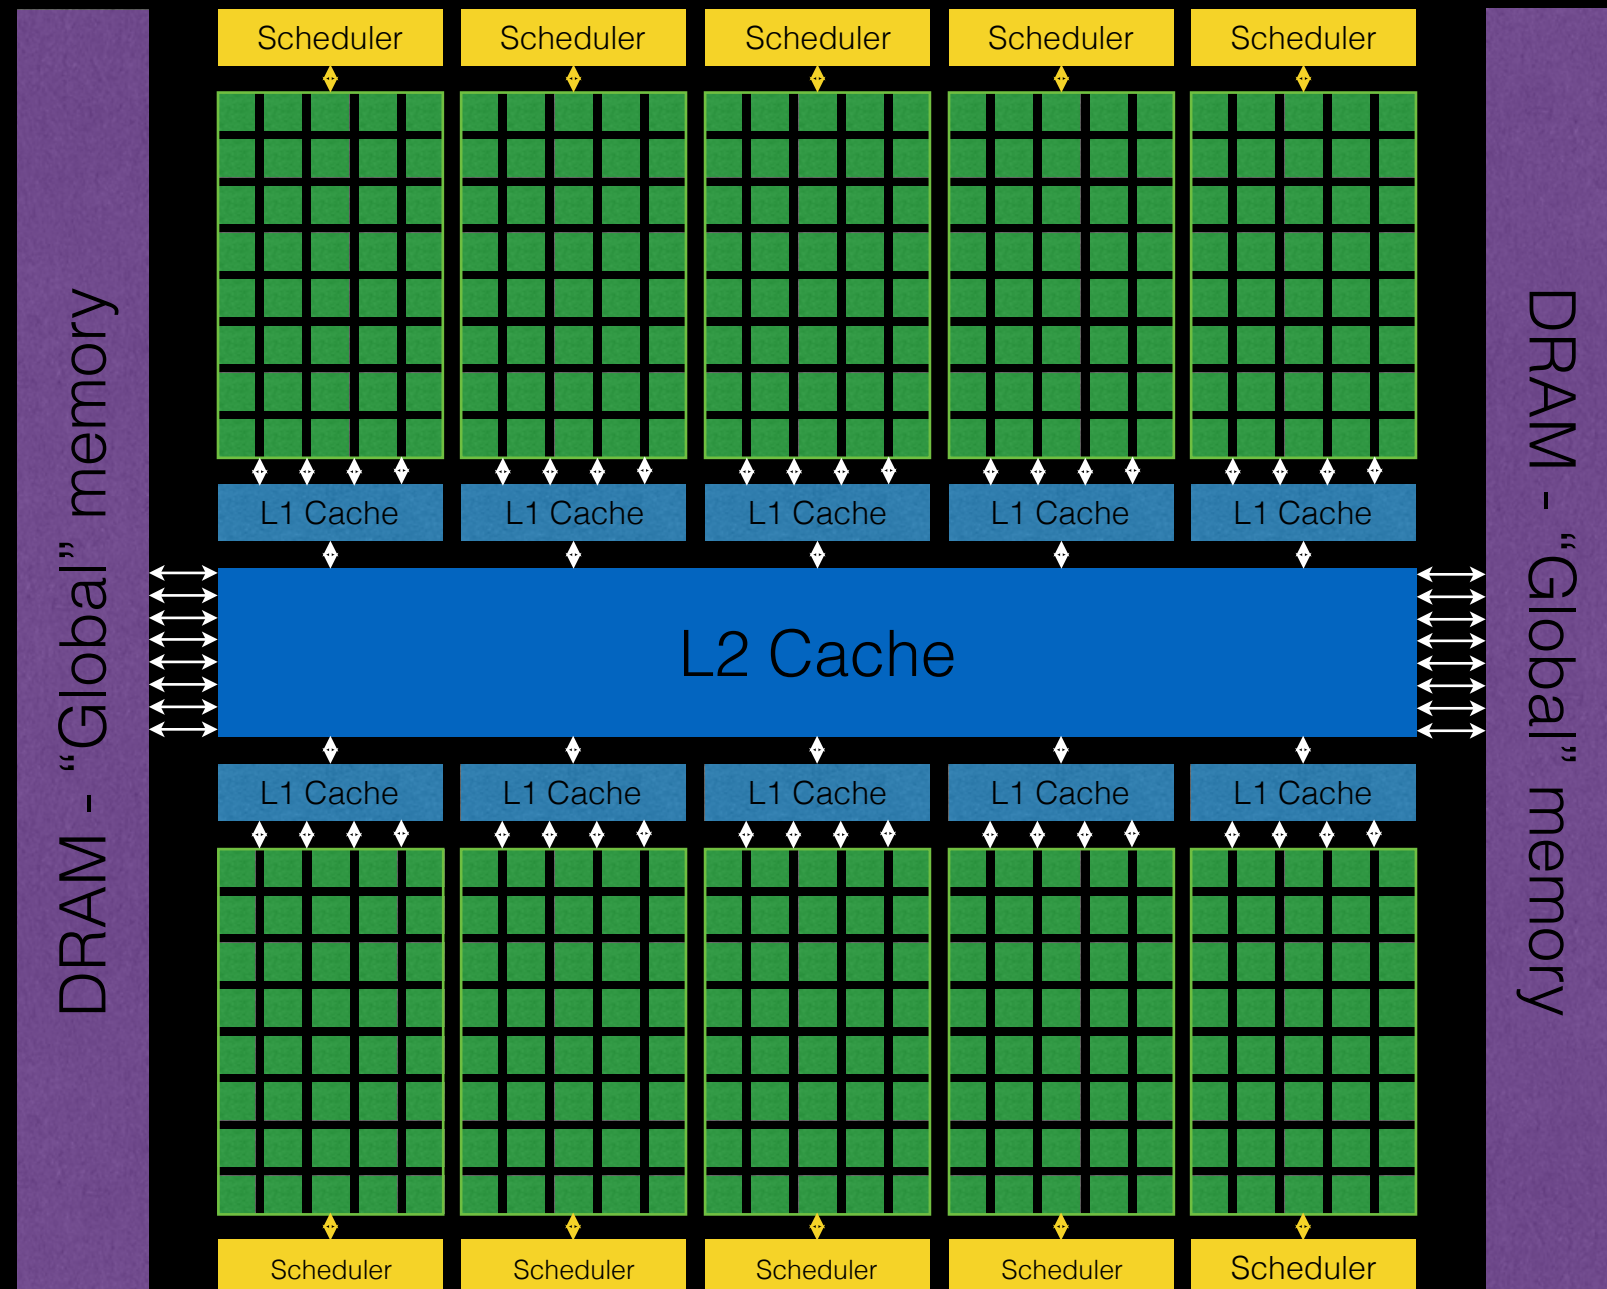
# Current CPUs



Context switching between threads is **expensive**:  
architecture optimised for single threaded performance




# Current GPUs



DRAM - "Global" memory

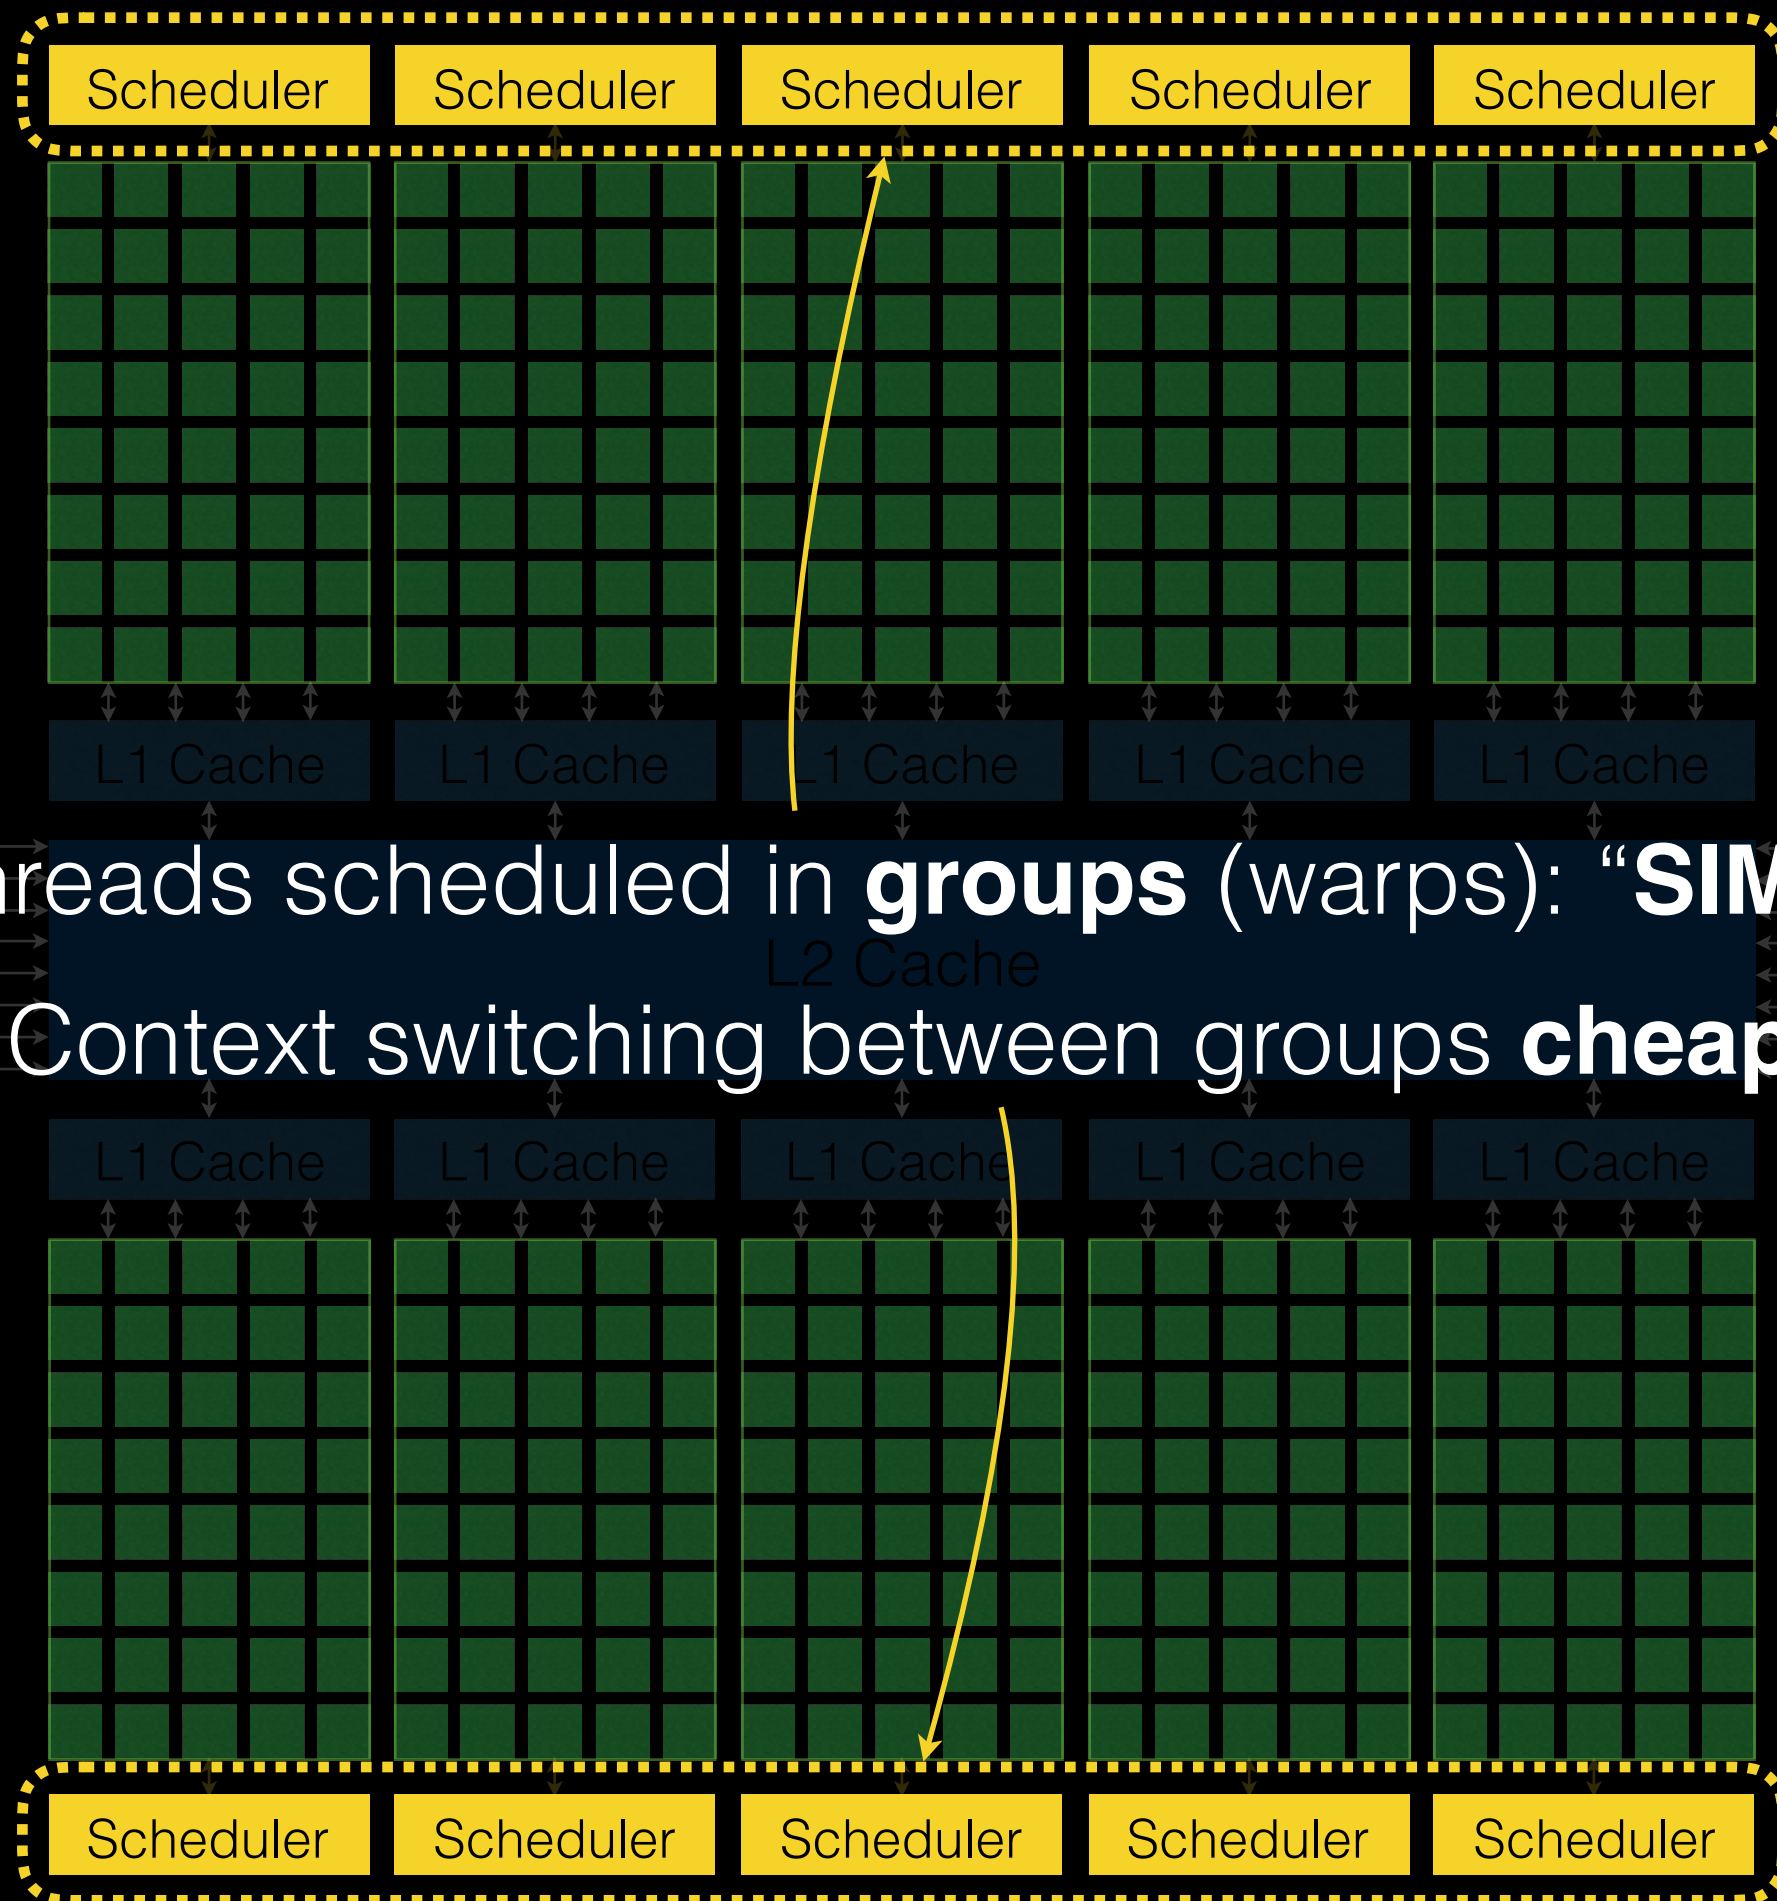
**Huge numbers of lightweight, simple cores**  
= massive parallelism

DRAM - "Global" memory





DRAM - "Global" memory



DRAM - "Global" memory

Threads scheduled in **groups** (warps): "**SIMT**"  
Context switching between groups **cheap**

**Conclusion:** GPUs give  
you *greater*, more *flexible*  
parallelism

**Conclusion:** GPUs give  
you *greater*, more *flexible*  
parallelism...right?

**Conclusion:** GPUs give  
you *greater*, more *flexible*  
parallelism...right?

**Why don't we use  
them for everything?**

**No caches:** Programmers must be careful with intelligently fetching data

**Grouped Parallelism:** Programmers should avoid thread divergence within groups

**Simple cores:** Programmers must parallelise with fine grained *and statically* distributed workloads



**No caches:** Programmers must be careful with intelligently fetching data

**Grouped Parallelism:** Programmers should avoid thread divergence within groups

**Simple cores:** Programmers must parallelise with fine grained *and statically* distributed workloads

GPUs encourage **static** *array based* parallelism

**No caches:** Programmers must be careful with intelligently fetching data

**Grouped Parallelism:** Programmers should avoid thread divergence within groups


**Simple cores:** Programmers must parallelise with fine grained *and statically* distributed workloads

GPUs encourage **static** *array based* parallelism  
How can we **efficiently** parallelise  
*other domains?*

# Domain of interest: **Graph algorithms**

Aim: accelerate  
processing using  
**GPU parallelism**


# What are the **problems?**



# What are the **problems?**

## **Sparsity:**


- Complex data structures
- Random data access patterns
- Low compute to data ratio



# What are the **problems?**

## **Irregularity:**

- Unbalanced workload within data sets
- Control flow divergence
- Coarse granularity of parallelism






# Our proposal:

Translate problem to simpler domain:  
(sparse) **Linear Algebra**


Express algorithm with  
**high level parallelism**  
(aka, functional languages)

**Explore** space of implementations to find  
**high performance** solution




Reduce (add, 0)    o    Map (mult)

```
__kernel void dotProduct(  
    __global const elem_t* a,  
    __global const elem_t* b,  
    __global const elem_t* res,  
    __global const int len){  
    int id = get_global_id(0);  
    elem_t tmp;  
    for(int i = id; i < len; i += get_global_size(0)){  
        tmp = a[i] * b[i];  
        atomic_add(res, tmp);  
    }  
}
```



# Our proposal:




$$\sum_{i=0} a_i \times b_i$$

Translate problem to simpler domain:  
(sparse) **Linear Algebra**

(intuition: operations on adjacency matrix of graph)

# Our proposal:



Implement linear algebra in parallel using a  
high level **functional** language

# Our proposal:


Use rewrite rules to automatically generate **high performance** OpenCL code

$$\sum_{i=0} a_i \times b_i$$

Reduce (add, 0)    ○    Map (mult)

```
__kernel void dotProduct(  
    __global const elem_t* a,  
    __global const elem_t* b,  
    __global const elem_t* res,  
    __global const int len){  
    int id = get_global_id(0);  
    elem_t tmp;  
    for(int i = id; i < len; i += get_global_size(0)){  
        tmp = a[i] * b[i];  
        atomic_add(res, tmp);  
    }  
}
```






**Sparsity:** Translate complex sparse domain to simpler dense domain

**Irregularity:** Regularise code using intelligent compilation techniques


---

**Current work:** extending to and evaluating on other domains.







AMD



NVIDIA



AMD



NVIDIA